


Generative Models

2025.07.10.

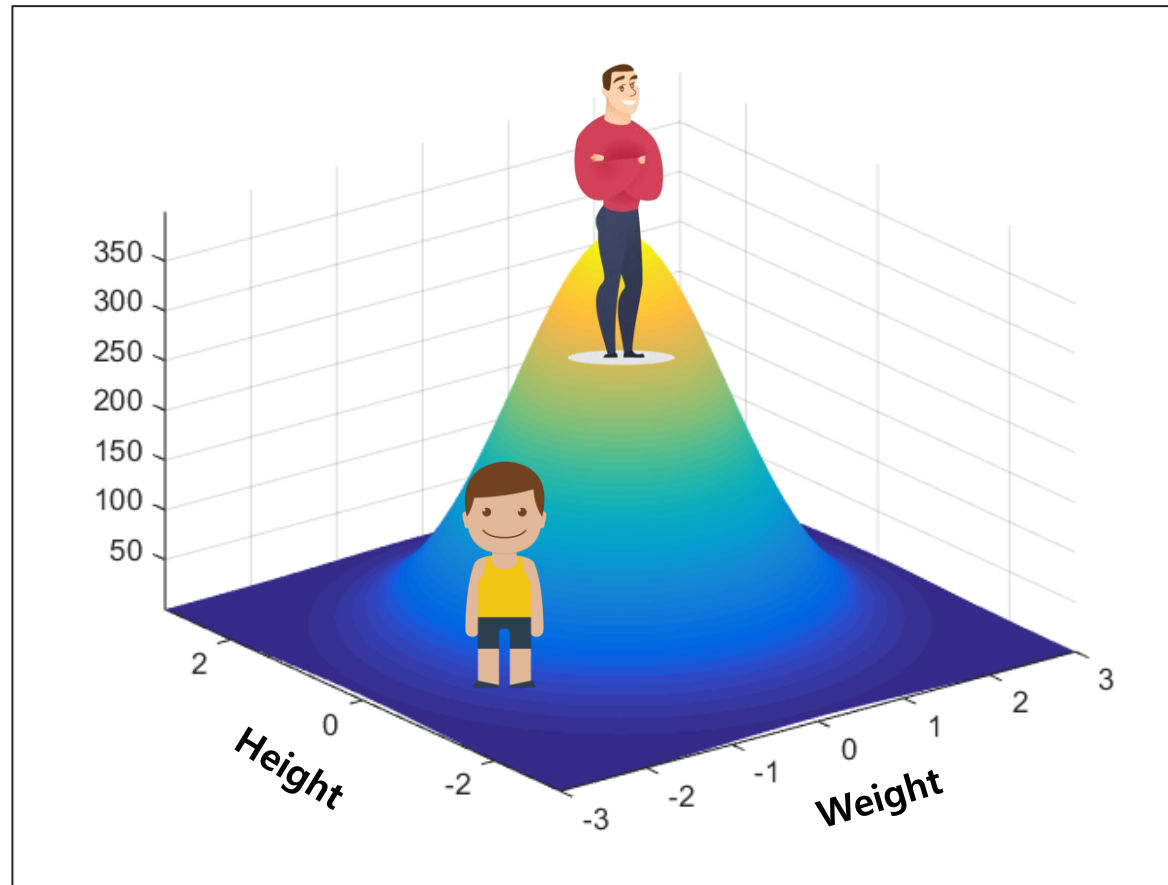
Copyright©2025 by 고재균



Generative Models [1]

- **Generative Models**



- Generative model aims to capture probability distribution of data, $P_{data}(x)$
- We assume that data comes from $P_{data}(x)$



Generative Models [2]

- **Generative Models**

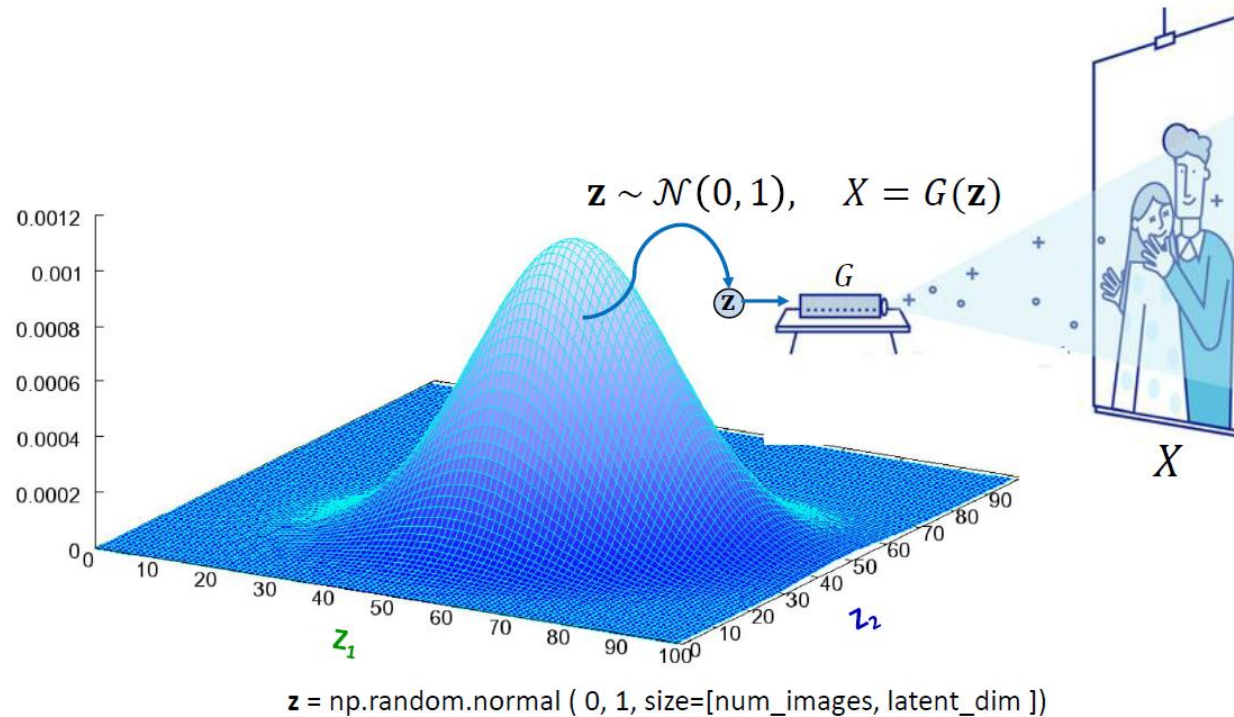
- Probability distribution of images?
- **Very complex!**



Generative Models [3]

- **Generative Models**

- What about making a simple distribution, and then projecting to a real-world complex distribution?
 - ▶ Making a simple distribution (e.g., standard normal) is easy!
 - ▶ Sample a **latent variable z** and then transform it to an image!



Generative Models [4]

- Generative Models

- Observed data x is originated from latent variables z



Myth of Cave

Generative Models [5]

- **Variational Autoencoder (VAE)**

- Assuming that training data \mathbf{x} is originated from underlying (unobserved) latent variables \mathbf{z}
- It defines intractable density function with latent \mathbf{z} :

$$P_{\theta}(\mathbf{x}) = \int P_{\theta}(\mathbf{x}|\mathbf{z}) P_{\theta}(\mathbf{z}) d\mathbf{z}$$

- We cannot optimize it directly because it is **intractable**!
- Instead, we will derive and maximize a **lower bound** on the **likelihood**


$$\begin{aligned} &= \underbrace{E_{\mathbf{z}}[\log P_{\theta}(\mathbf{x})] - D_{KL}\left(q_{\phi}(\mathbf{z}|\mathbf{x}) || P_{\theta}(\mathbf{z})\right)}_{\text{Tractable lower bound}} + \underbrace{D_{KL}\left(q_{\phi}(\mathbf{z}|\mathbf{x}) || P_{\theta}(\mathbf{z}|\mathbf{x})\right)}_{\substack{\text{Interactable} \\ \text{(KL-Divergence } \geq 0)}} \\ &\geq E_{\mathbf{z}}[\log P_{\theta}(\mathbf{x})] - D_{KL}\left(q_{\phi}(\mathbf{z}|\mathbf{x}) || P_{\theta}(\mathbf{z})\right) \end{aligned}$$

Generative Models [6]

- **Variational Autoencoder (VAE)**

- Revisiting Autoencoder (AE)

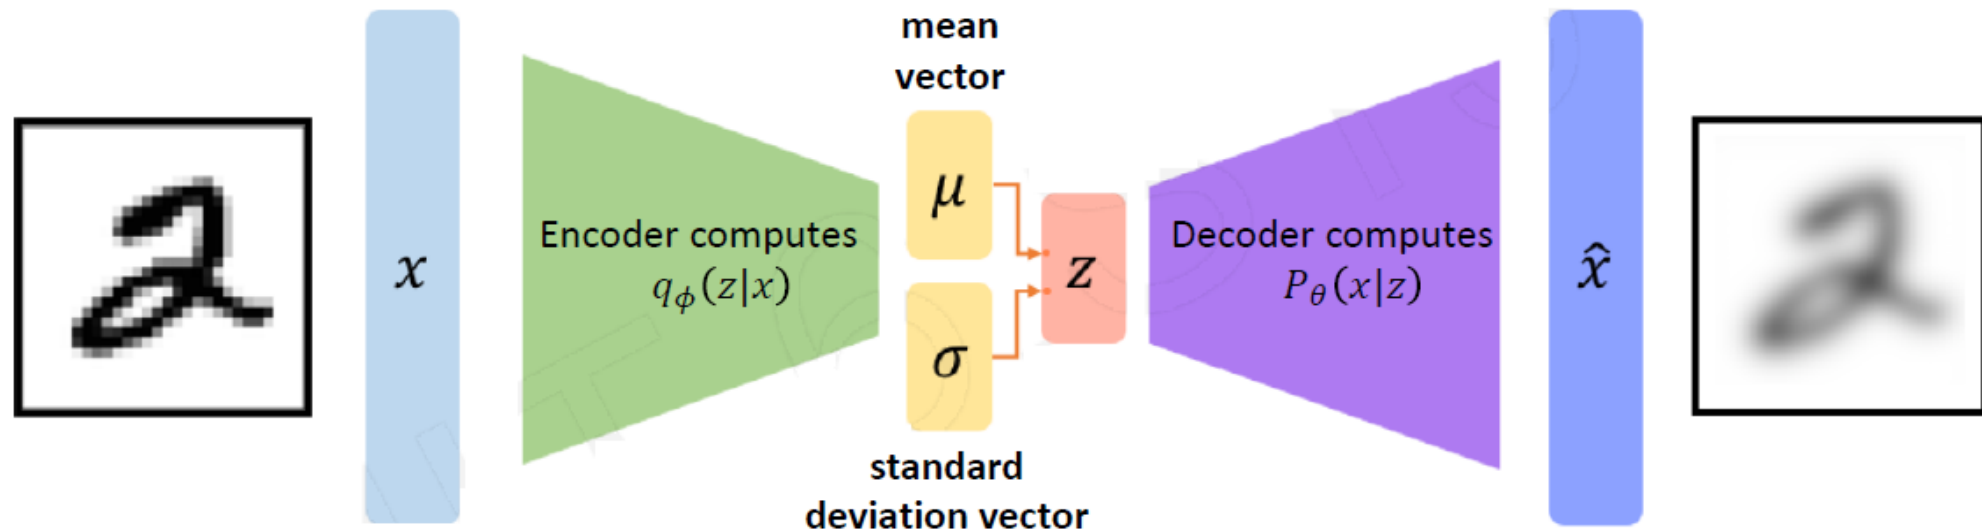
- ▶ Unsupervised approach for learning a **lower-dimensional feature** representation from training data (without explicit labels)
 - ▶ Train such that latent features can be used to **reconstruct** the original data ("Autoencoding" - encoding itself)
 - ▶ Goal is to learn the latent features \mathbf{z} that capture (or encode) as much information about the data \mathbf{x} as possible



Generative Models [7]

- **Variational Autoencoder (VAE)**

- Assume training data x is generated from underlying (unobserved) **latent variables** z
- **Encoder network** models $q_{\phi}(z|x)$ and **decoder network** models $P_{\theta}(x|z)$



Sample z from $z|x \sim \mathcal{N}(\mu_{z|x}, \Sigma_{z|x})$

Generative Models [8]

- **Variational Autoencoder (VAE)**

- Maximizing **log-likelihood** of training data x : $\log P_{\theta}(x)$

$$\log P_{\theta}(x) = E_{z \sim q_{\phi}(z|x)} [\log P_{\theta}(x)] \quad \text{(Taking expectation with regards to } z \text{)}$$

$$= E_z \left[\log \frac{P_{\theta}(x|z) P_{\theta}(z)}{P_{\theta}(z|x)} \right] \quad \text{(Bayes' rule : } \frac{P(B|A)P(A)}{P(B)} \text{)}$$

$$= E_z \left[\log \frac{P_{\theta}(x|z) P_{\theta}(z)}{P_{\theta}(z|x)} \frac{q_{\phi}(z|x)}{q_{\phi}(z|x)} \right] \quad \text{(Multiply & divide with the same term)}$$

$$= E_z [\log P_{\theta}(x|z)] - E_z \left[\log \frac{q_{\phi}(z|x)}{P_{\theta}(z)} \right] + E_z \left[\log \frac{q_{\phi}(z|x)}{P_{\theta}(z|x)} \right]$$

$$= \boxed{E_z [\log P_{\theta}(x|z)] - D_{KL} (q_{\phi}(z|x) || P_{\theta}(z))} + \boxed{D_{KL} (q_{\phi}(z|x) || P_{\theta}(z|x))}$$

Tractable lower bound

Interactable



(KL-Divergence ≥ 0)

$$\geq E_z [\log P_{\theta}(x|z)] - D_{KL} (q_{\phi}(z|x) || P_{\theta}(z))$$

Generative Models [9]

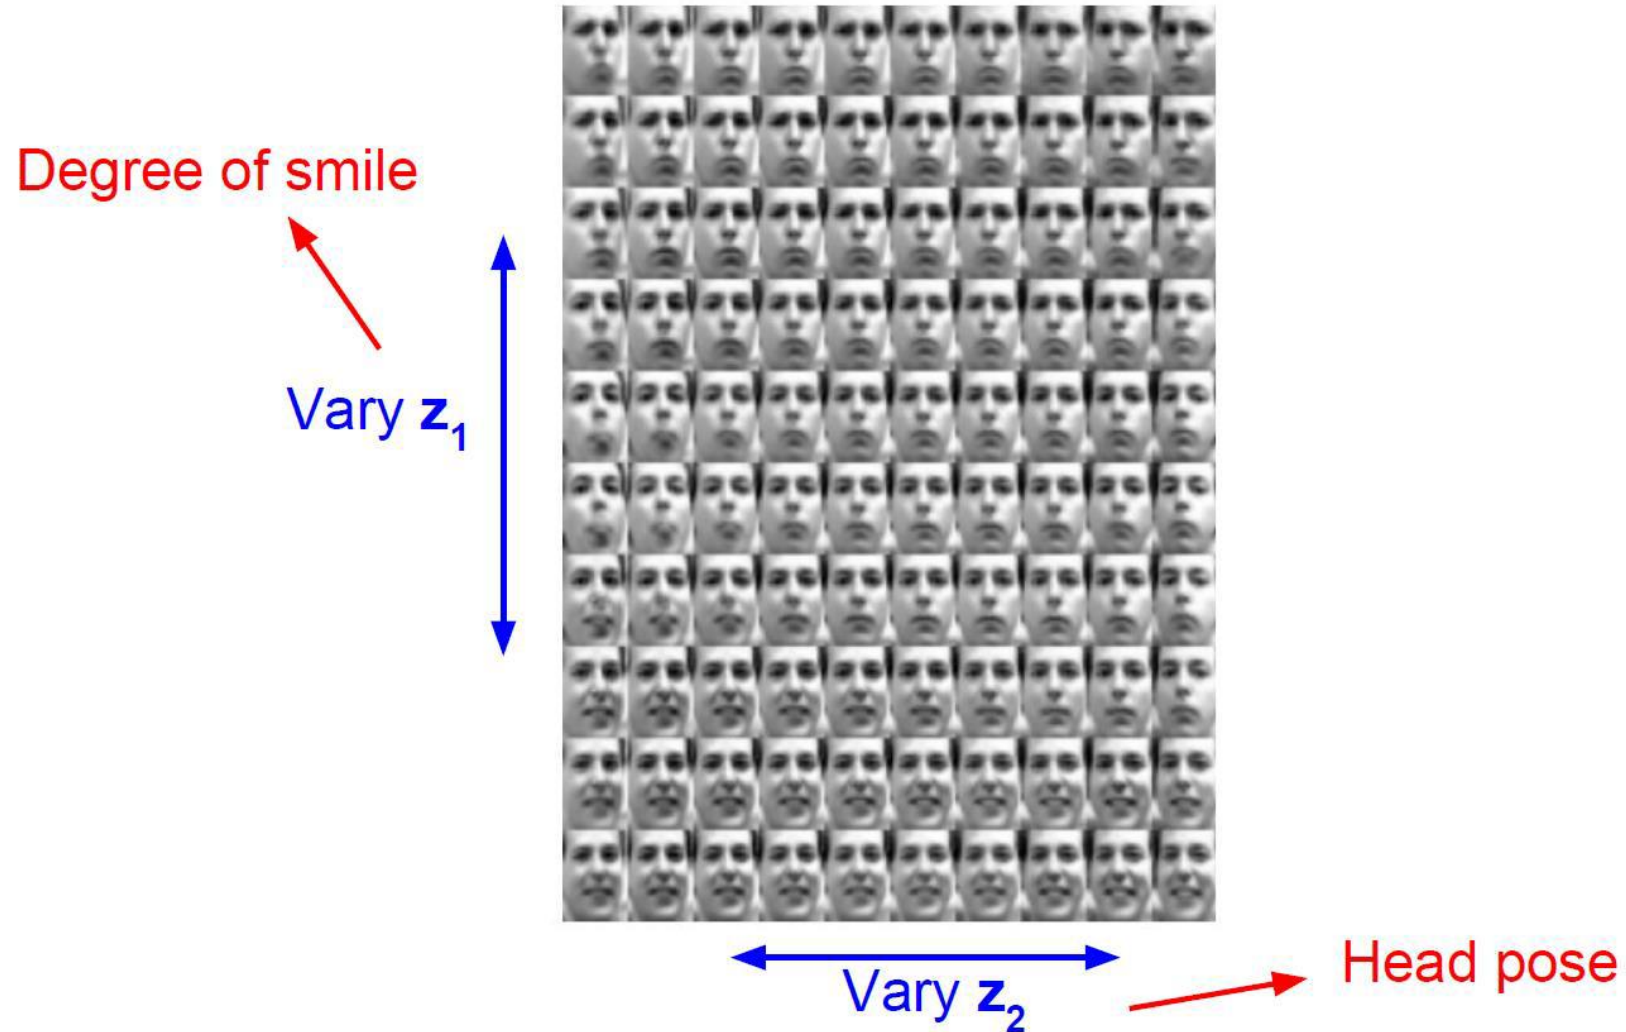
• Variational Autoencoder (VAE)

- Generates samples with regarding to \mathbf{z}
- Only decoder is used here, and \mathbf{z} is produced and varied manually
- Slowly increase or decrease a single latent variable while keeping all other variables fixed
- Each number is smoothly transitioning to another number



Generative Models [10]

- Variational Autoencoder (VAE)



Generative Models [11]

- Variational Autoencoder (VAE) 실습

Variational Autoencoder를 활용한 MNIST 데이터셋 생성

(Google Colab. 환경)


Generative Models [12]

- **Variational Autoencoder (VAE)**

- Defines an intractable density function
 - ▶ It derives and optimizes a lower bound on likelihood of training data instead

- **Generative Adversarial Network (GAN)**

- Gives up on explicitly modeling density function, but just wants ability to generate data



Generative Models [13]

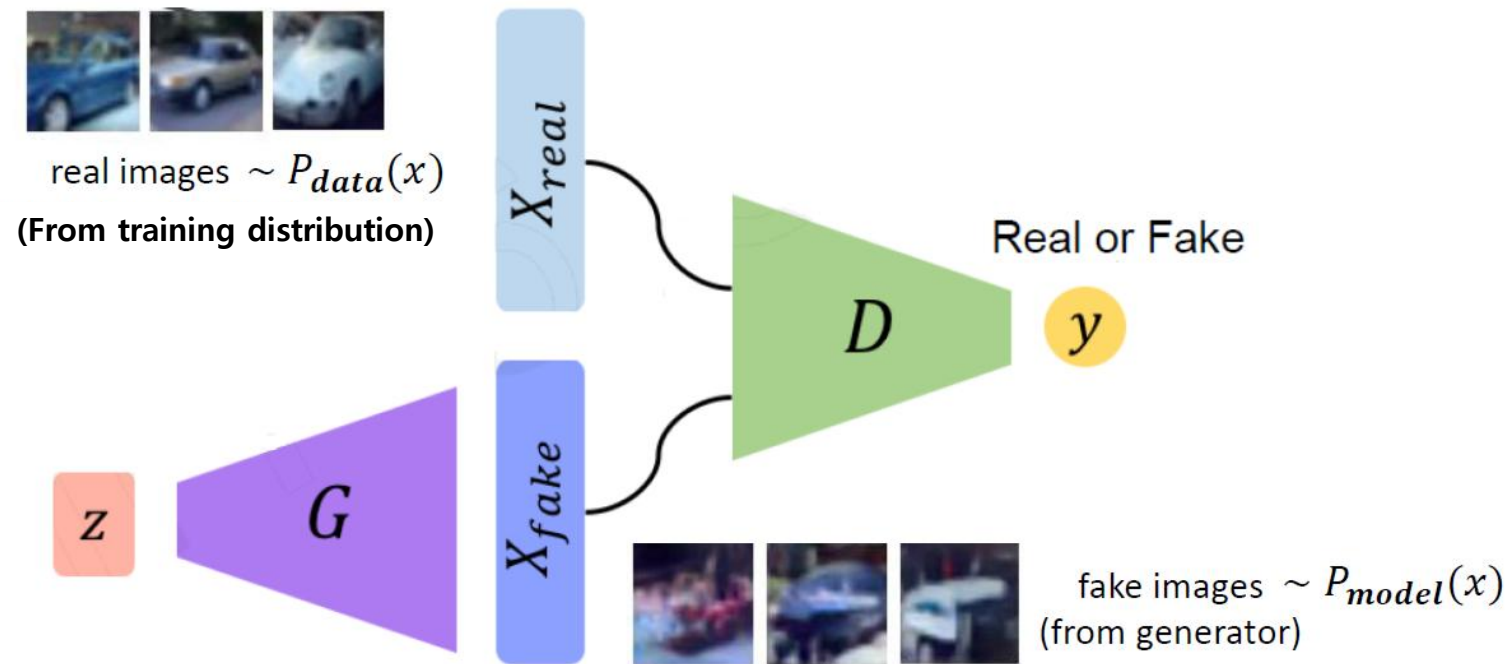
- **Generative Adversarial Network (GAN)**

- Gives up explicitly defining and estimating a probability
- Just wants ability to generate data from training distribution $P_{data}(x)$
- Approach
 - ▶ Sample latent variable z from just a **simple distribution**, e.g., normal Gaussian $\mathcal{N}(\mathbf{0}, \mathbf{1})$ (a.k.a. random noise)
 - ▶ Then, learn a mapping function (**generator**) from z to training distribution $P_{data}(x)$
- Training Strategy
 - ▶ 2-player game strategy (경찰과 도둑)
 - ▶ Employ another model, named **discriminator** guiding **generator** to training distribution $P_{data}(x)$
 - ▶ Make the generator and the discriminator compete with each other
 - ▷ Generator tries to **fool** the **discriminator** by generating **real-looking data**
 - ▷ Discriminator tries to **distinguish** between **real** and **fake data**

Generative Models [14]

- **Generative Adversarial Network (GAN)**

- To succeed in this game, the generator must learn to generate data that is indistinguishable from real-world data
- Hence, needs to generate data that **looks** drawn from the same distribution as the training data



Generative Models [15]

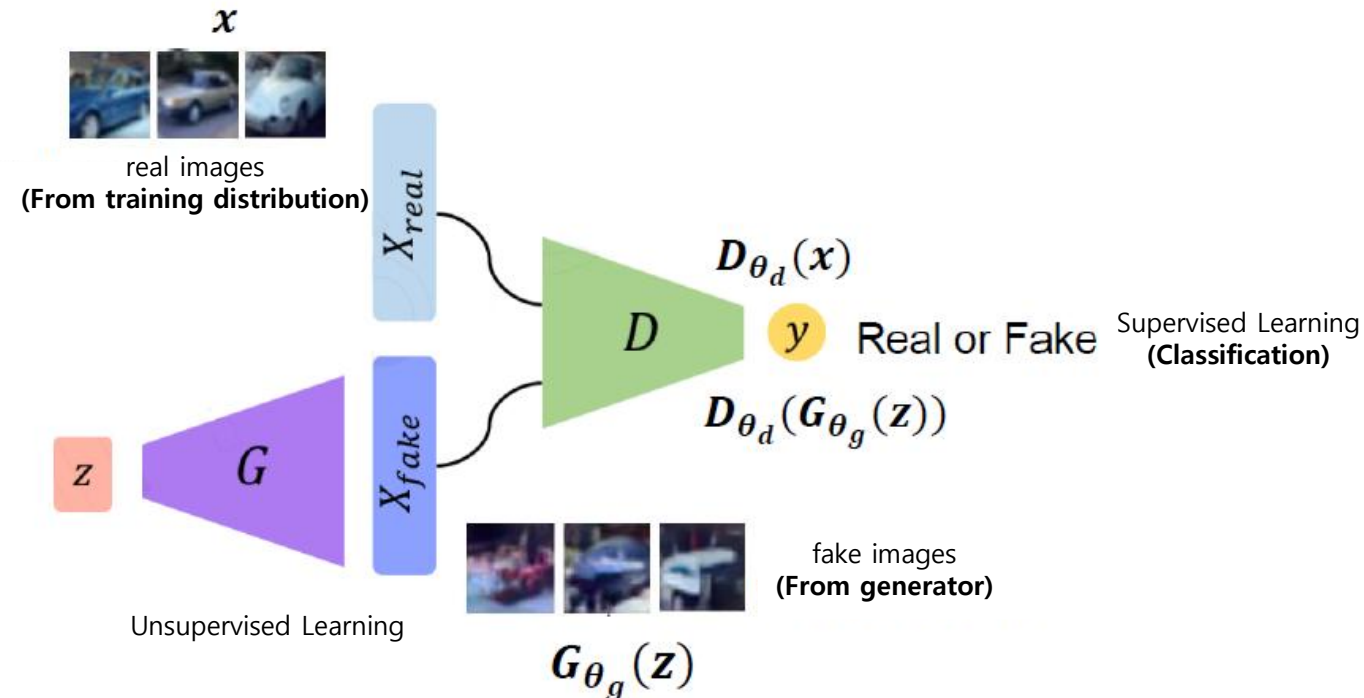
• Generative Adversarial Network (GAN)

▪ Notations

► $D_{\theta_d}(x)$ → Discriminator's output : likelihood that x is a real data, range of $[0, 1]$

► $G_{\theta_g}(z)$ → Generated fake data

► $D_{\theta_d}(G_{\theta_g}(z))$ → Likelihood that $G_{\theta_g}(z)$ is a real data, range of $[0, 1]$



Generative Models [16]

- Generative Adversarial Network (GAN)

- Notations

- ▶ $D_{\theta_d}(x) \rightarrow$ Discriminator's output : likelihood that x is a real data, range of $[0, 1]$

- ▶ $G_{\theta_g}(z) \rightarrow$ Generated fake data

- ▶ $D_{\theta_d}(G_{\theta_g}(z)) \rightarrow$ Likelihood that $G_{\theta_g}(z)$ is a real data, range of $[0, 1]$

- Objective Function

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

- ▶ Discriminator wants to maximize objective such that $D_{\theta_d}(x)$ is close to 1 (real) and $D_{\theta_d}(G_{\theta_g}(z))$ is close to 0 (fake)

- ▶ Generator wants to minimize objective such that $D_{\theta_d}(G_{\theta_g}(z))$ is close to 1

- (discriminator is fooled into thinking generated $G_{\theta_g}(z)$ is real)

Generative Models [17]

- **Generative Adversarial Network (GAN)**

- Alternate between:

- ▶ Gradient **ascent** on discriminator, generator is fixed (not updated) in this step

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- ▶ Gradient **descent** on generator, discriminator is fixed in this step

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

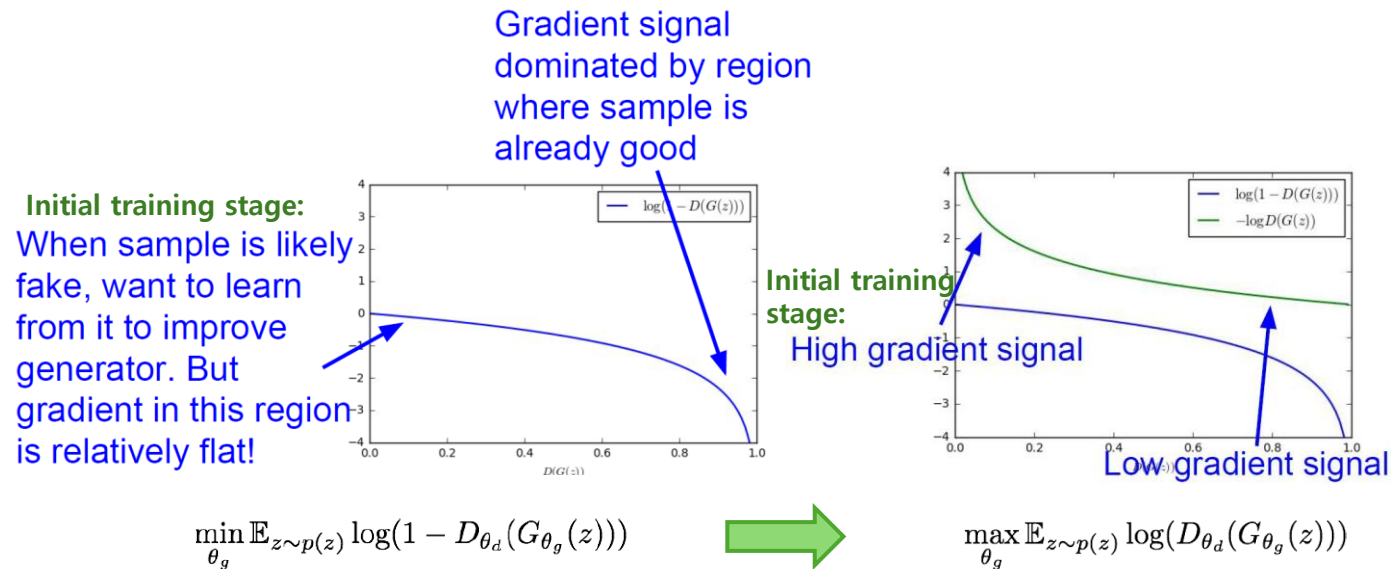
- ▶ In practice, however, optimizing this generator objective does **not** work well!
 - ▶ Hence, bottom objective function is widely-used!

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

Generative Models [18]

• Generative Adversarial Network (GAN)

▪ Training GAN



- ▶ Instead of minimizing likelihood of discriminator being correct, now maximize likelihood of discriminator being wrong
- ▶ Same objective of fooling discriminator, but now **higher gradient signal for initial training stage**
- ▶ It works much better! Standard in practice!

Generative Models [19]

- Generative Adversarial Network (GAN)

- Training GAN

- ▶ Learning algorithm

for number of training iterations **do**
 for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)}))) \right]$$

end for

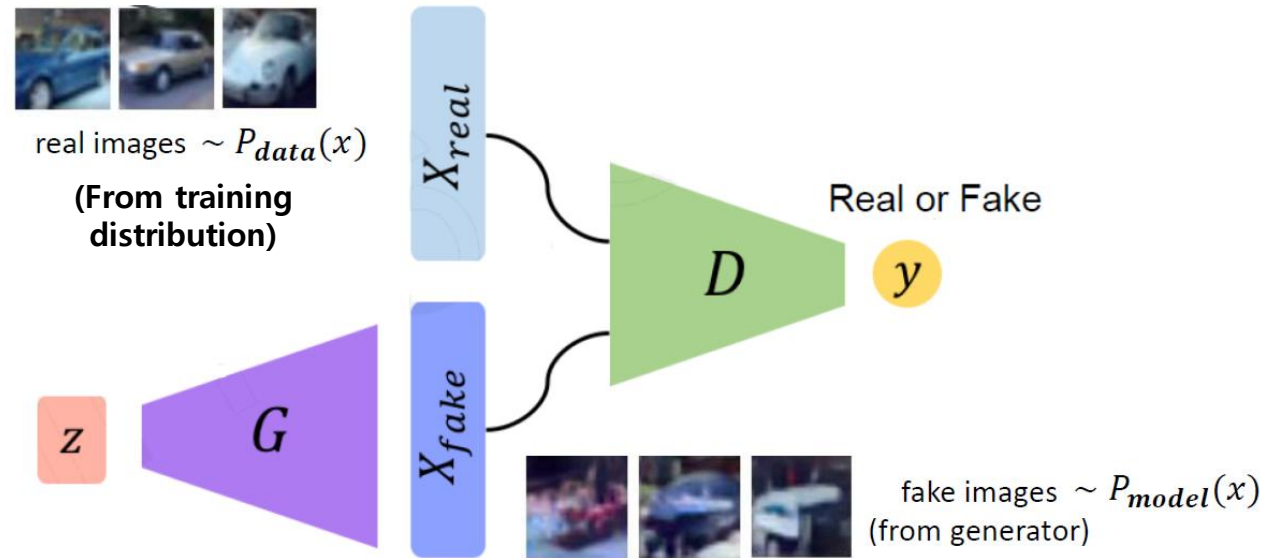
- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by ascending its stochastic gradient (improved objective):

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$

end for



Generative Models [20]

- Generative Adversarial Network (GAN)



Generative Models [21]

- Generative Adversarial Network (GAN)




Generative Models [22]

- Generative Adversarial Network (GAN)



Noise $\sim N(0,1)$





Generator
Network



Generative Models [23]

- Generative Adversarial Network (GAN)

- Conditional GAN



Class	0	1	2	...	9
One-hot vector	$\begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$	$\begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$...	$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{pmatrix}$

Generative Models [24]

- Generative Adversarial Network (GAN)

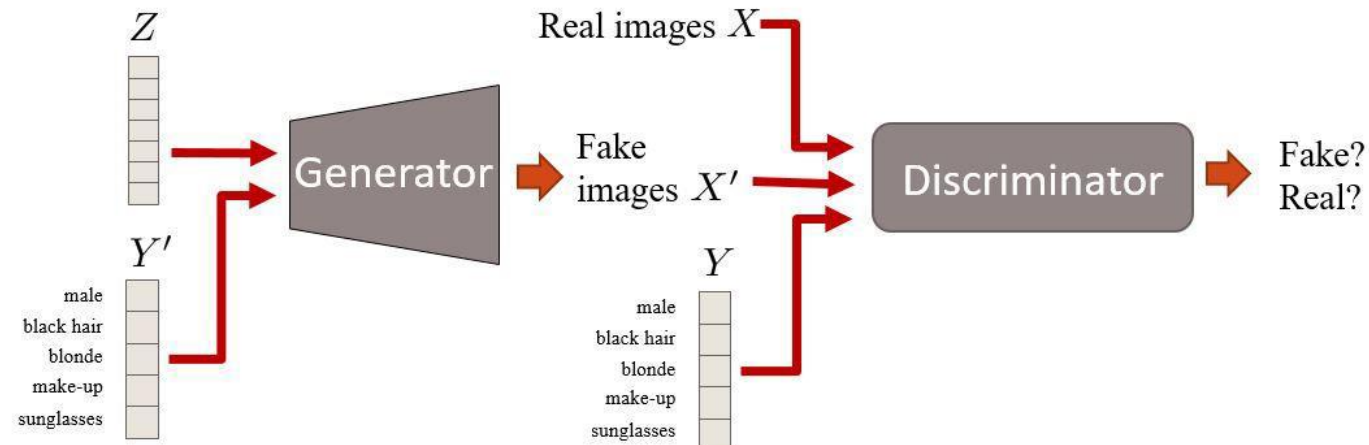
- Conditional GAN

- ▶ Training objective

$$\min_G \max_D \left(\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim p_{\text{data}}(\mathbf{x}, \mathbf{y})} [\log D(\mathbf{x}, \mathbf{y})] \right. \\ \left. + \mathbb{E}_{\mathbf{y} \sim p_{\mathbf{y}}, \mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z}, \mathbf{y}), \mathbf{y}))] \right)$$

- ▶ Implementation

- ▷ Just concatenate the input vectors of **G** and **D** (i.e., **z** and **x**) with the condition vector **y**



Generative Models [25]

• Generative Adversarial Network (GAN)

▪ More GANs



► Pix2Pix

▷ Image-to-Image (I2I) Translation


$$G^* = \arg \min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G).$$

$$\mathcal{L}_{GAN}(G, D) = \mathbb{E}_y[\log D(y)] + \mathbb{E}_{x,z}[\log(1 - D(G(x, z)))].$$

$$\mathcal{L}_{L1}(G) = \mathbb{E}_{x,y,z}[\|y - G(x, z)\|_1].$$




Positive examples



G tries to synthesize fake images that fool **D**

D tries to identify the fakes

Negative examples



Generative Models [26]

- Generative Adversarial Network (GAN)

- More GANs

- ▶ Pix2Pix

- ▷ Image-to-Image (I2I) Translation



Generative Models [27]

• Generative Adversarial Network (GAN)

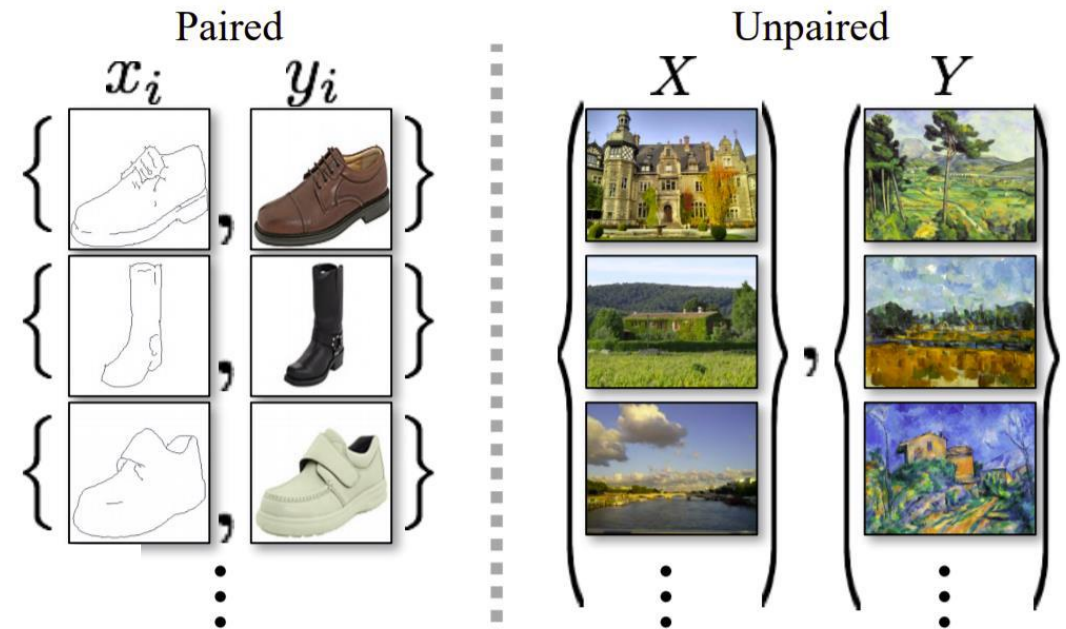


▪ More GANs

► DiscoGAN & CycleGAN

▷ Image-to-Image (I2I) Translation (Paired & Unpaired)

$$G^*, F^* = \arg \min_{G, F} \max_{D_X, D_Y} \mathcal{L}(G, F, D_X, D_Y)$$

$$\begin{aligned} \mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F), \end{aligned}$$





Generative Models [28]

- Generative Adversarial Network (GAN)

- More GANs

- ▶ DiscoGAN & CycleGAN

- ▷ Image-to-Image (I2I) Translation (Paired & Unpaired)



horse → zebra



zebra → horse

Generative Models [29]

- Generative Adversarial Network (GAN)

- More GANs

- ▶ DiscoGAN & CycleGAN

- ▷ Image-to-Image (I2I) Translation (Paired & Unpaired)



winter Yosemite → summer Yosemite



summer Yosemite → winter Yosemite



apple → orange



orange → apple



Generative Models [30]

- Generative Adversarial Network (GAN)


- More GANs

- StarGAN

(a) Cross-domain models



(b) StarGAN

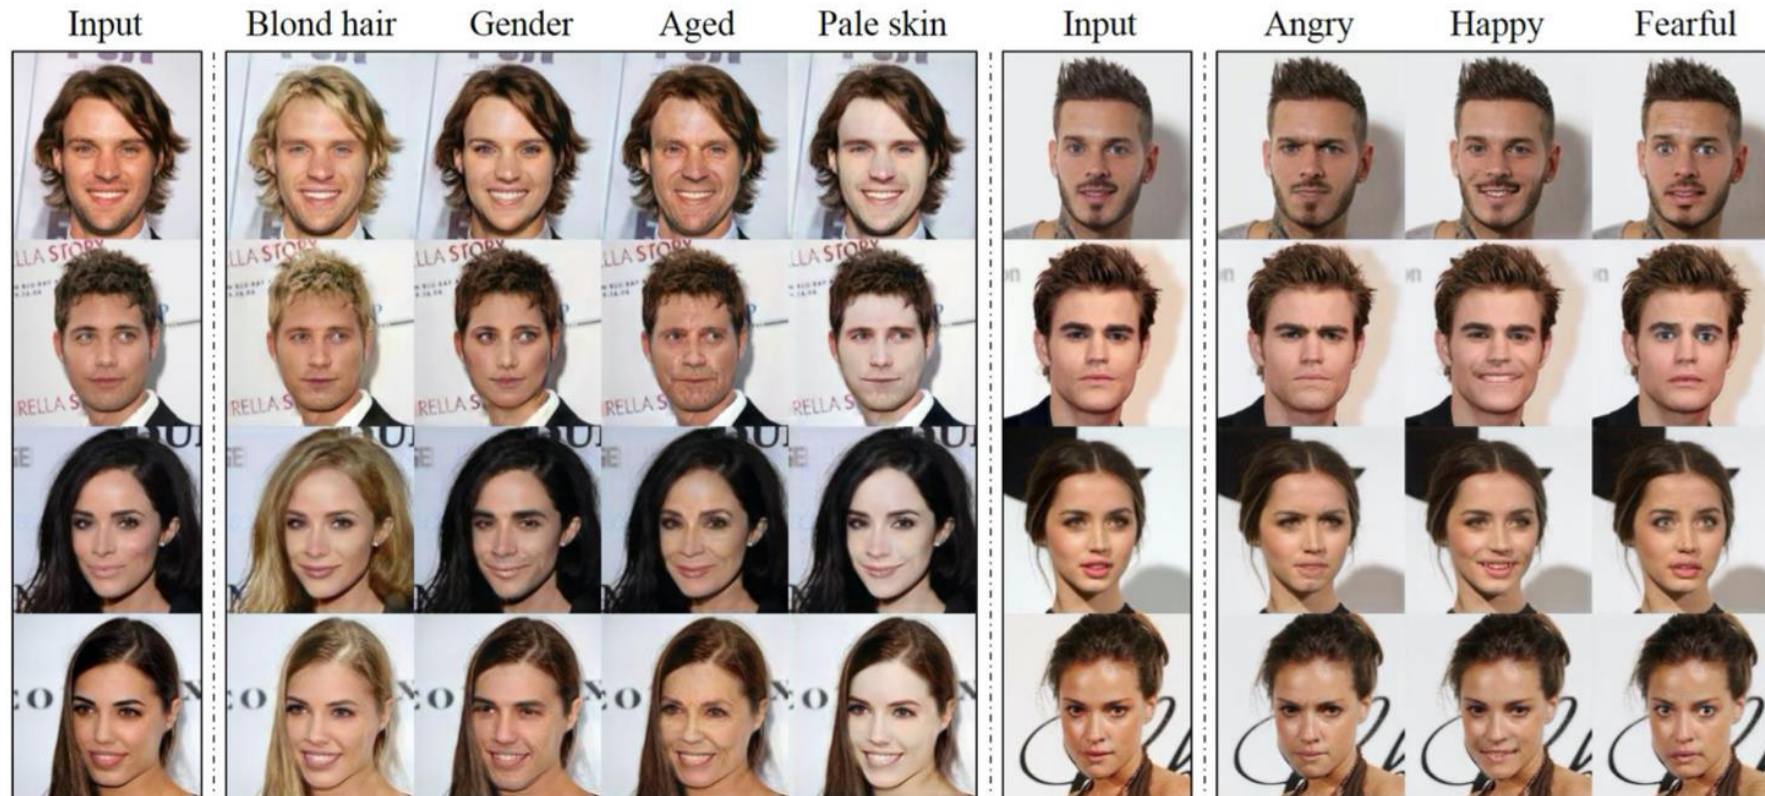


Generative Models [31]

- Generative Adversarial Network (GAN)

- More GANs

- StarGAN




Generative Models [32]

- **Generative Adversarial Network (GAN)**

- More GANs

- ▶ Image Inpainting




Generative Models [33]

- **Generative Adversarial Network (GAN)**

- More GANs

- ▶ Category-to-Image



Generative Models [34]

- **Generative Adversarial Network (GAN)**

- More GANs

- ▶ Category-to-Image



(a) Snail



(b) Studio couch



(c) Harvester

Generative Models [35]

- **Generative Adversarial Network (GAN)**

- More GANs

- ▶ Caption-to-Image (Text-to-Image)



a red car parked on the side of a road

a blue car parked on the side of a road



a pizza on a plate at a restaurant

someone is just about to cut the pizza



oranges on a table next to a liquor bottle

a pile of oranges sitting in a wooden crate

Generative Models [36]

- Generative Adversarial Network (GAN) 실습

Generative Adversarial Network를 활용한 Fashion MNIST데이터셋 생성

(Google Colab. 환경)